

chap1-hw-mlops-zoomcamp-2023

May 19, 2023

1 MLOps-zoomcamp-2023-chap1-intro-hw

Use the “Run” button to execute the code.

1.1 Download the Dataset

```
[1]: dataset_link_1 = "https://d37ci6vzurychx.cloudfront.net/trip-data/
      ↪yellow_tripdata_2022-01.parquet"
      dataset_link_2 = "https://d37ci6vzurychx.cloudfront.net/trip-data/
      ↪yellow_tripdata_2022-02.parquet"
```

```
[2]: !pwd
      !wget $dataset_link_1
      !wget $dataset_link_2
```

```
/content
--2023-05-19 14:07:36-- https://d37ci6vzurychx.cloudfront.net/trip-
data/yellow_tripdata_2022-01.parquet
Resolving d37ci6vzurychx.cloudfront.net (d37ci6vzurychx.cloudfront.net)...
65.8.245.178, 65.8.245.171, 65.8.245.50, ...
Connecting to d37ci6vzurychx.cloudfront.net
(d37ci6vzurychx.cloudfront.net)|65.8.245.178|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 38139949 (36M) [application/x-www-form-urlencoded]
Saving to: 'yellow_tripdata_2022-01.parquet'

yellow_tripdata_202 100%[=====>] 36.37M 59.2MB/s in 0.6s

2023-05-19 14:07:37 (59.2 MB/s) - 'yellow_tripdata_2022-01.parquet' saved
[38139949/38139949]

--2023-05-19 14:07:37-- https://d37ci6vzurychx.cloudfront.net/trip-
data/yellow_tripdata_2022-02.parquet
Resolving d37ci6vzurychx.cloudfront.net (d37ci6vzurychx.cloudfront.net)...
65.8.245.178, 65.8.245.171, 65.8.245.50, ...
Connecting to d37ci6vzurychx.cloudfront.net
(d37ci6vzurychx.cloudfront.net)|65.8.245.178|:443... connected.
HTTP request sent, awaiting response... 200 OK
```

Length: 45616512 (44M) [application/x-www-form-urlencoded]
Saving to: 'yellow_tripdata_2022-02.parquet'

yellow_tripdata_202 100%[=====>] 43.50M 69.9MB/s in 0.6s

2023-05-19 14:07:37 (69.9 MB/s) - 'yellow_tripdata_2022-02.parquet' saved
[45616512/45616512]

```
[94]: import pandas as pd

import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.feature_extraction import DictVectorizer
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn.linear_model import Ridge
```

```
[95]: # Specify the file path of the parquet file
parquet_file_train = 'yellow_tripdata_2022-01.parquet'
parquet_file_val = 'yellow_tripdata_2022-02.parquet'

# Read the parquet file using pandas
df_train = pd.read_parquet(parquet_file_train)
```

1.2 Data Analysis and Exploration

```
[96]: df_train.shape
```

```
[96]: (2463931, 19)
```

1.2.1 Read the data for January. How many columns are there? 19

1.2.2 What's the standard deviation of the trips duration in January?

```
[97]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2463931 entries, 0 to 2463930
Data columns (total 19 columns):
 #   Column                Dtype
---  -
 0   VendorID              int64
 1   tpep_pickup_datetime  datetime64[ns]
```

```

2  tpep_dropoff_datetime  datetime64[ns]
3  passenger_count        float64
4  trip_distance          float64
5  RatecodeID             float64
6  store_and_fwd_flag     object
7  PULocationID           int64
8  DOLocationID           int64
9  payment_type           int64
10 fare_amount            float64
11 extra                  float64
12 mta_tax                float64
13 tip_amount             float64
14 tolls_amount           float64
15 improvement_surcharge  float64
16 total_amount           float64
17 congestion_surcharge   float64
18 airport_fee            float64
dtypes: datetime64[ns](2), float64(12), int64(4), object(1)
memory usage: 357.2+ MB

```

```

[98]: df_train['duration'] = df_train.tpep_dropoff_datetime - df_train.
      ↪ tpep_pickup_datetime
      df_train.duration = df_train.duration.apply(lambda td: td.total_seconds() / 60)

```

```

[99]: sns.distplot(df_train['duration'], label='duration')
      plt.legend()

```

<ipython-input-99-16a96132d705>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_train['duration'], label='duration')
```

```
[99]: <matplotlib.legend.Legend at 0x7f68d041b9d0>
```



```
[100]: # Calculate the standard deviation of the column
std_dev = df_train['duration'].std()

# Print the result
print('Standard Deviation:', std_dev)
```

Standard Deviation: 46.44530513776802


What's the standard deviation of the trips duration in January? 46.445 ~ 46.45

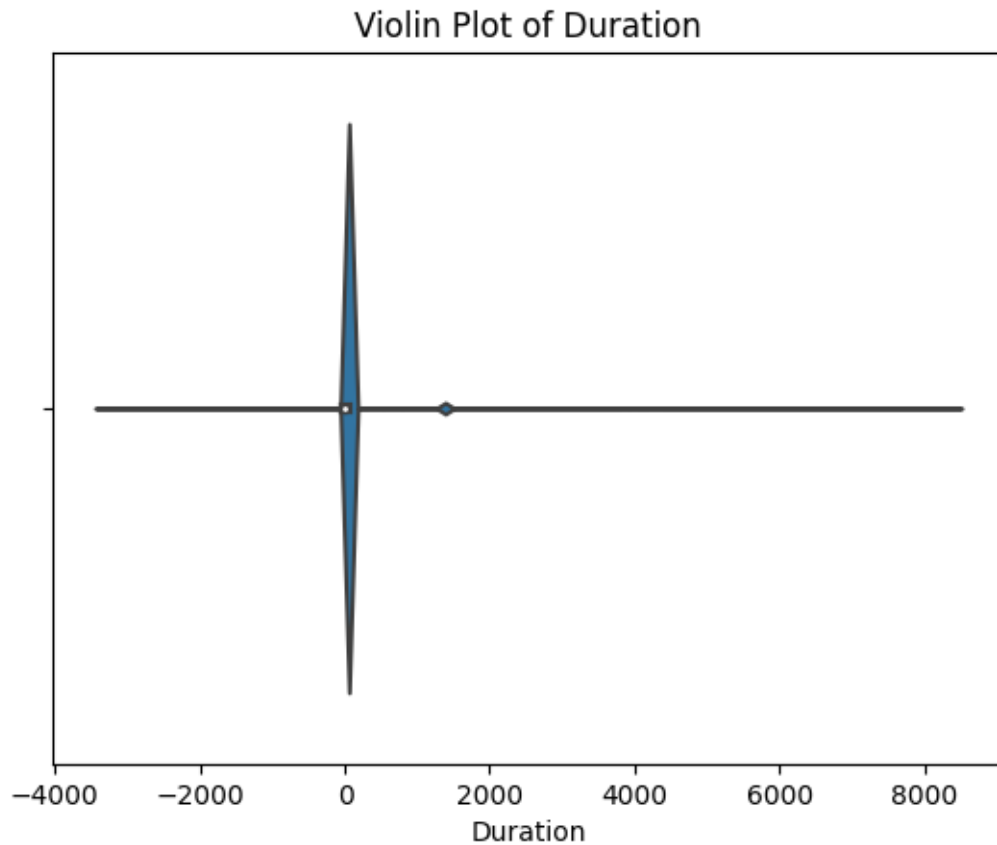
1.2.3 What fraction of the records left after you dropped the outliers?

```
[101]: # Boxplot
sns.boxplot(x=df_train['duration'])
plt.title('Boxplot of Duration')
plt.xlabel('Duration')
plt.show()

# Violin plot
sns.violinplot(x=df_train['duration'])
plt.title('Violin Plot of Duration')
plt.xlabel('Duration')
```

```
plt.show()
```






```
[102]: df_train = df_train[(df_train.duration >= 1) & (df_train.duration <= 60)]
```


```
[103]: df_train.shape
```

```
[103]: (2421440, 20)
```

```
[104]: # Boxplot
sns.boxplot(x=df_train['duration'])
plt.title('Boxplot of Duration')
plt.xlabel('Duration')
plt.show()

# Violin plot
sns.violinplot(x=df_train['duration'])
plt.title('Violin Plot of Duration')
plt.xlabel('Duration')
plt.show()
```





Still the values could be trimmed by shrinking the upper limit from 60 to 30. But as the questions states to trim between 1 and 60 so we have to stick to the same.

```
[105]: print(f"The fraction of records left after dropping the outliers is_
↪ {round(df_train.shape[0]*100/2463931,2)} %")
```

The fraction of records left after dropping the outliers is 98.28 %

1.2.4 One-hot encoding

```
[106]: df_train.columns
```

```
[106]: Index(['VendorID', 'tpep_pickup_datetime', 'tpep_dropoff_datetime',
            'passenger_count', 'trip_distance', 'RatecodeID', 'store_and_fwd_flag',
            'PULocationID', 'DOLocationID', 'payment_type', 'fare_amount', 'extra',
            'mta_tax', 'tip_amount', 'tolls_amount', 'improvement_surcharge',
            'total_amount', 'congestion_surcharge', 'airport_fee', 'duration'],
            dtype='object')
```

```
[107]: # selected_columns = ['PULocationID', 'DOLocationID']
```



```
[108]: categorical = ['PULocationID', 'DOLocationID']
       numerical = []
```

```
[109]: def read_dataframe(filename):
       if filename.endswith('.parquet'):
           df = pd.read_parquet(filename)

       df['duration'] = df.tpep_dropoff_datetime - df.tpep_pickup_datetime
       df.duration = df.duration.apply(lambda td: td.total_seconds() / 60)

       df = df[(df.duration >= 1) & (df.duration <= 60)]

       categorical = ['PULocationID', 'DOLocationID']
       df[categorical] = df[categorical].astype(str)

       return df
```

```
[110]: df_val = read_dataframe(parquet_file_val)
       df_val.shape
```

```
[110]: (2918187, 20)
```

```
[111]: len(df_train), len(df_val)
```

```
[111]: (2421440, 2918187)
```

```
[112]: train_dicts = df_train[categorical + numerical].to_dict(orient='records')
```

```
[113]: dv = DictVectorizer()
       X_train = dv.fit_transform(train_dicts)
```

```
[114]: X_train.shape, type(X_train)
```

```
[114]: ((2421440, 2), scipy.sparse._csr.csr_matrix)
```

```
[115]: print(X_train[0])
```

```
(0, 0)      236.0
(0, 1)      142.0
```

```
[116]: feature_names = dv.get_feature_names_out()
       dimensionality = len(feature_names)
       print(f'the dimensionality of this matrix (number of columns):_
       ↪{dimensionality}')
```

```
the dimensionality of this matrix (number of columns): 2
```

1.3 Training a model

```
[117]: target = 'duration'
       y_train = df_train[target].values
```

```
[118]: lr = LinearRegression()
       lr.fit(X_train, y_train)
```

```
[118]: LinearRegression()
```

```
[119]: from sklearn.metrics import mean_squared_error
```

```
[120]: y_pred = lr.predict(X_train)
```

```
[121]: print(f'the RMSE of the model on the training data {mean_squared_error(y_train, y_pred, squared=False)}')
```

the RMSE of the model on the training data 8.920327827581444

```
[122]: # Plotting the distributions
       def plot_distribution(actual, predicted):
           sns.distplot(actual, hist=True, kde=False, bins=10, label='Actual Values')
           sns.distplot(predicted, hist=True, kde=False, bins=10, label='Predicted Values')
           plt.xlabel('Values')
           plt.ylabel('Frequency')
           plt.legend()
           plt.title('Distribution of Predicted and Actual Values')
           plt.show()
```

```
[123]: plot_distribution(y_train, y_pred)
```

<ipython-input-122-fc33e34a5dbe>:3: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(actual, hist=True, kde=False, bins=10, label='Actual Values')
<ipython-input-122-fc33e34a5dbe>:4: UserWarning:
```


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with

similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(predicted, hist=True, kde=False, bins=10, label='Predicted Values')
```



1.4 Evaluating the model

```
[124]: val_dicts = df_val[categorical + numerical].to_dict(orient='records')
X_val = dv.transform(val_dicts)
```

```
[125]: y_val = df_val[target].values
```

```
[126]: y_pred = lr.predict(X_val)
```

```
[128]: print(f'the RMSE of the model on the training data {mean_squared_error(y_val, y_pred, squared=False)}')
```

the RMSE of the model on the training data 10.032124095033593

```
[129]: plot_distribution(y_val,y_pred)
```

```
<ipython-input-122-fc33e34a5dbe>:3: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(actual, hist=True, kde=False, bins=10, label='Actual Values')
<ipython-input-122-fc33e34a5dbe>:4: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(predicted, hist=True, kde=False, bins=10, label='Predicted
Values')
```



```
[130]: y_pred
```

```
[130]: array([16.56034382, 16.56034382, 16.56034382, ..., 16.56034382,
        16.56034382, 16.56034382])
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]: # df_train['PU_DO'] = df_train['PULocationID'] + '_' + df_train['DOLocationID']
     # df_val['PU_DO'] = df_val['PULocationID'] + '_' + df_val['DOLocationID']
```

1.5 Similar things for the validation Dataset

```
[ ]: # df_val['duration'] = df_val.tpep_dropoff_datetime - df_val.  
      ↳ tpep_pickup_datetime  
      # df_val.duration = df_val.duration.apply(lambda td: td.total_seconds() / 60)  
  
      # # Boxplot  
      # sns.boxplot(x=df_val['duration'])  
      # plt.title('Boxplot of Duration')  
      # plt.xlabel('Duration')  
      # plt.show()  
  
      # # Violin plot  
      # sns.violinplot(x=df_val['duration'])  
      # plt.title('Violin Plot of Duration')  
      # plt.xlabel('Duration')  
      # plt.show()
```

```
[90]: from sklearn.feature_extraction import DictVectorizer  
  
      # Create an instance of DictVectorizer  
      vectorizer = DictVectorizer()  
  
      # Fit the vectorizer to your data  
      data = [{'feature1': 2, 'feature2': 5}, {'feature1': 1, 'feature3': 3}]  
      vectorizer.fit(data)  
  
      # Get the feature names  
      feature_names = vectorizer.get_feature_names_out()  
  
      # Determine the dimensionality  
      dimensionality = len(feature_names)  
  
      # Print the dimensionality  
      print("Dimensionality:", dimensionality)
```

Dimensionality: 3

```
[91]: feature_names
```

```
[91]: array(['feature1', 'feature2', 'feature3'], dtype=object)
```

```
[ ]:
```